Work without Claude

Choose one exercise. Keep a 20-minute timer. Talk, test the idea, and fill in the worksheet. Use the prepared case below or an industry you know. The student keeps time and takes notes.

Prepared case: a small advisory firm

Fictional case. Every operational number below is invented for the workshop. This is business-research work, not investment, legal, or clinical advice.

A six-person firm prepares public-information market briefs for business clients. Each brief combines research, a draft, and senior review. The firm wonders whether agents could help it provide a useful service to smaller clients over the next three to five years.

Starting question: Could the same team offer more briefs without reducing quality? How would that change its service, customers, or competition?

Starting expectation to accept or replace: "Faster drafting will let us serve more clients." Record the table's own view first.

Today: a deliberately small model

- Five analysts each have 24 hours per week for these briefs: **120 analyst hours** total.
- One senior reviewer has **12 hours per week** for review.
- Each brief takes **8 analyst hours**, then **1 review hour**.
- There is demand for **20 briefs per week**.
- Work is divisible for capacity planning; round completed weekly briefs down to whole numbers.

Proposed change: assumptions to test

- Agent-assisted research and drafting reduce analyst time to **4 hours per brief**.
- Extra checking raises senior review to **1.5 hours per brief**.
- Staff and demand initially stay the same.
- Agent cost, training, implementation, queueing, and client acquisition are outside the simple model. Discuss them before claiming business value.

Three things the table may change

1. Analyst hours per brief with agents.
2. Senior review hours per brief with agents.
3. Weekly demand.

The student can run python3 scripts/simulate.py or use materials/simulation-worked-example.md on paper.

Evidence cards: use even without web access

Prepared October 15, 2026. These are short paraphrases of original sources, not full papers. If offline, cite the original URL and say "prepared summary; not rechecked live." The empirical examples describe different tasks and older tools. They are not estimates of autonomous agents' performance in 2026 or proof about the fictional advisory firm.

E1: benefits can depend on the worker and task

Source: Brynjolfsson, Li, and Raymond, *Generative AI at Work*, NBER Working Paper 31161, original 2023 working-paper version. Original version; paper record.

Finding: The original study of a conversational assistant in customer support reported roughly 14% more issues resolved per hour on average, with larger benefits for less experienced workers.

Boundary: One workplace and task setting; assisted support work, not an autonomous advisory service. Later paper versions may report revised estimates. Do not turn 14% into a universal productivity assumption.

Ask: Which parts of our work resemble this setting? Which do not?

E2: perceived speed can differ from measured results

Source: METR, *Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity*, July 10, 2025. Study report.

Finding: In a randomized study of 16 experienced developers completing 246 tasks in familiar repositories, allowing early-2025 AI tools increased completion time by 19%. Developers' impressions were more positive than the measured result.

Boundary: Specific developers, repositories, tasks, and tools. This does not establish that AI slows all work, or describe today's tools.

Newer context: In its February 24, 2026 update, METR reported that selection and measurement problems made its follow-up an unreliable estimate of current effects. It considered greater speedups plausible but could not reliably quantify them.

Ask: What would we measure rather than relying on how fast the work feels?

E3: a skill can package a repeatable process

Source: Anthropic, *Equipping agents for the real world with Agent Skills*, October 16, 2025; updated December 18, 2025. Engineering explanation.

Observation: Skills package instructions and optional resources that an agent can load for a task.

Boundary: This is a vendor's description of a capability, not an independent evaluation of business outcomes. A skill cannot make its own claims true.

Ask: Which parts of our team's expertise could be written down and reused? Which require situated judgment?

Use the cards well

E1 and E2 concern different settings: their results cannot be averaged or treated as a head-to-head comparison. E3 establishes a design pattern, not a return on investment. Together they motivate testing the whole workflow. If your industry is different, say that direct industry evidence remains missing. For a current-industry claim, find an appropriate primary source live or leave the claim open.

Research: is it happening?

- **0-2:** Choose one claim. Say what you expect. Example: "Agents could help our firm serve more clients."

- **2-7:** Read evidence cards E1 and E2. For each, name what it found and what it cannot tell us about our firm.

- **7-12:** Ask: Which study is more relevant to our task? What important evidence is missing? Do the differing results necessarily contradict each other?

- **12-16:** Revise the claim. Say what might change for the service or customers. What must people still do? What changed in your thinking?

- **16-20:** Fill in the six-part worksheet. Cite the cards' original sources and mark them "prepared summaries; not rechecked live." Preserve disagreement.

If you use another industry, the cards may be analogies only. "We need industry evidence" is a useful finding. This is a short evidence scan, not a comprehensive review.

Finish with the six sections and a succinct one-paragraph summary. Name your table, then ask the student to save and push its reviewed report using docs/report-out.md.

Simulation: what changes?

- **0-2:** Ask whether the firm could finish more briefs. Say what you expect.

- **2-6:** Read the case and worked numbers. Challenge one assumption.

- **6-11:** Compare 12 briefs today with eight under the agent assumptions. Pick one input to change. Recalculate using the smaller of analyst capacity, review capacity, and demand; round down. Use the worked table if you prefer.

- **11-16:** Ask what limits completed work. What costs, quality issues, or customer needs are missing? What might change in the service or competitive landscape? What changed in your thinking?

- **16-20:** Fill in the worksheet. Label the numbers fictional. Include the changed assumption and a measurement that would test it.

Changing only agent review time to 0.75 hours gives 16 completed briefs in the prepared model. This does not prove that quality stays equal or that clients will buy them.

Finish with the six sections and a succinct one-paragraph summary. Name your table, then ask the student to save and push its reviewed report using docs/report-out.md.

Role playing: who sees it differently?

- **0-2:** Decide whether the firm should pilot a brief service for smaller clients. Say your initial view.

- **2-5:** Choose three or four role cards. Assign roles to willing participants.

- **5-10:** Each person gives one need and one worry in about 20 seconds. Which conflict matters most?

- **10-14:** Add this fictional condition: clients require traceable sources for every important claim. Revise the pilot, or explain why you would decline it.

- **14-16:** Step out of character. What did we assume about these people? Who would we need to ask? What changed in our thinking?

- **16-20:** Fill in the worksheet. Include the conflict, revised idea, and one real stakeholder question.

Imagined reactions are hypotheses. They are not survey findings or evidence that customers would buy the service.

Finish with the six sections and a succinct one-paragraph summary. Name your table, then ask the student to save and push its reviewed report using docs/report-out.md.

Forecasting: what could make change faster or slower?

- **0-2:** Pick a concrete change by 2030. Say what you currently expect.

- **2-6:** Choose one uncertain driver: review quality, data access, customer trust, or implementation capacity.

- **6-11:** Describe a faster-change future and a slower-change future. For each: what happens to the work, what must be true, and what people still contribute?

- **11-14:** Name two signs you could actually observe. Say what to measure, where to get it, and when to check. Example: completed briefs at unchanged quality during a documented pilot.

- **14-16:** Pick an action useful in both futures. What changed in your thinking?

- **16-20:** Fill in the worksheet. Include both futures, the separating assumption, and the signs to watch.

These are possible paths, not calibrated predictions. Any proposed threshold is an assumption to discuss, not an industry benchmark.

Finish with the six sections and a succinct one-paragraph summary. Name your table, then ask the student to save and push its reviewed report using docs/report-out.md.

Problem solving: what could stop this from working?

- **0-2:** Define success. Say what you expect. Example: more completed briefs at equal quality.

- **2-6:** Pick the main obstacle. In the prepared model, review is limiting; the reason for slow review remains untested.

- **6-10:** Compare three responses: improve source packets, narrow the service, or keep the current process. Choose one to test.

- **10-14:** Pick one constraint card. Which response survives? What needs to change?

- **14-16:** Design a small test: responsible role, scope, comparison, result to measure, quality check, and stop/continue rule. What changed in your thinking?

- **16-20:** Fill in the worksheet. Record the response, constraint, test, and uncertainty.

This is a test plan. You have not yet shown that the response works. Do not treat a plausible explanation as a proven root cause.

Finish with the six sections and a succinct one-paragraph summary. Name your table, then ask the student to save and push its reviewed report using docs/report-out.md.

Simulation you can do on paper

All inputs and outputs are fictional capacity calculations, not forecasts.

Analyst capacity = 120 analyst hours / analyst hours per brief. Review capacity = 12 reviewer hours / review hours per brief. Completed briefs = round DOWN the smallest of analyst capacity, review capacity, and demand.

Case	Analyst hours / brief	Review hours / brief	Analyst capacity	Review capacity	Demand	Completed / week
Today	8	1	15	12	20	12
Agent assumption	4	1.5	30	8	20	8
Change only review time to 1	4	1	30	12	20	12
Change only review time to 0.75	4	0.75	30	16	20	16
Agent assumption, demand falls to 6	4	1.5	30	8	6	6

Example: faster drafting doubles analyst capacity, but the assumed checking burden reduces completed briefs from 12 to 8. This illustrates a possible bottleneck; it does not demonstrate that agents cause more checking in real firms.

Ask: Is 1.5 hours realistic? What could make review easier? What is outside the model? If more briefs become feasible, would clients want them? What happens when competitors can do the same?

This is a steady weekly capacity bound. It omits arrivals, queues, variability, rework loops, quality, money, and implementation costs. It cannot establish profit, staffing needs, or actual service levels. Changing review time does not prove equal quality: test that separately.

Role cards

Use for a fictional decision: **Should the advisory firm pilot an agent-assisted brief service for smaller clients?** These are invented interests, not descriptions of real people or evidence of public opinion. People may change the cards. Take roles by interest, not demographic identity.

Role	What I want	What worries me	A question I can ask
Customer	Useful advice I can act on	Paying for generic work or hidden errors	What can I verify, and who answers when it is wrong?
Analyst	Better work and time to learn	More output targets without more support	What happens to the time saved and to my development?
Senior reviewer	Reliable work and a sustainable workload	Errors that take longer to find than to avoid	What reaches my desk, and what evidence comes with it?
Firm leader	A valuable, viable service	Costs, reputation, and weak demand	What must we measure before expanding?
Competitor	A reason clients choose us	Everyone offering the same cheap brief	If everyone has this tool, what is distinctive?

Use three or four roles. Include a person affected by the change and a person accountable for results. If the industry requires an oversight perspective, substitute a generic oversight role and ask what rules would need checking; do not invent laws.

Round 1: each person states one need and one objection, in 20 seconds. Claude may play **one missing role** in two sentences. Round 2: change a condition and revise the decision. Debrief out of character: what did we assume about these people, and who should we actually ask?

Add one constraint

Invented prompts, not actual events. Let the table choose one, or use the first.

1. **No more reviewer time.** You cannot hire or add overtime for six months.
2. **Higher checking burden.** A client's audit requires a traceable source for every material claim.
3. **Demand doubles.** More clients want the service, but the team size is unchanged.
4. **Less data access.** The agent can use public information only.
5. **Less customer trust.** A prospective client will pay only after seeing a verified sample.
6. **Competitors catch up.** Rival firms can buy the same agent tools.

Ask: "Which response still works, what must change, and what would we test first?" A constraint is part of the thought experiment; it is not new evidence.